


The Activation Protocol

Get a new user to a real win on day one, with the agent doing the work.



What's inside:

- The Level Design model: why activation is Level 1 of a game, and how the agent runs it.
- How to define the one first-session win that matters, and what to do when your product can't deliver value on day one.
- The path, the first quest, and the win moment, built so the agent walks the user through all of it.

Contents

–	Introduction	03
–	The game: Level Design for software	04
–	Summary cheat sheet	05
01	Why users leave before they start	06
02	Define the same-day win	07
03	Proof beats value	08
04	Clear the path, let the agent walk it	09
05	The scripted first quest	10
06	Land the win, start the loop	11
–	The Activation Protocol	12
	The fill-in worksheet	
–	Run this with your favorite LLM	14

The protocol for a product that onboards itself

START HERE

I've watched the same thing happen for twenty years. A team builds something genuinely good, pours a year into it, launches, and then watches people sign up, poke around for a few minutes, and never come back. The product was rarely the problem. The users left before they ever saw what it could do.

I'm Bal. I've spent about twenty years across marketing and product in B2B SaaS, working with teams at companies like Google, Microsoft, and Intel. For most of that time I was designing onboarding without calling it that. Every product I ever touched had one, the stretch between "just signed up" and "this is part of how I work now," and I treated it as part of product design, because that's what it is. Onboarding is that whole stretch, not the sign-up wizard. The wizard is the first few minutes; the real job runs from the first win, to the habit that brings them back, to the day they bring their team in. I've seen a handful done well. I've seen most done badly, and almost always in the same way.

Here's the pattern in the bad ones. The team knows activation is weak, so they reach for a tool. Appcues, Pendo, a setup checklist, a row of tooltips. They bolt a layer of guidance on top of the product and call it onboarding. The activation rate barely moves, because the tooltips were never the problem. The product still drops a brand-new user into a blank screen and asks them to figure out what to do. A costume went on; nothing underneath it changed.

This protocol is about the opposite move, and it's written for where software is going: agent-led products, where there's an AI inside the product that can talk to the user and actually do things for them. In that world you don't hand a new user a checklist and wish them luck. The agent walks them in. It asks what they came for, builds their first thing with them, and narrates what's happening so they learn by watching it work. The user should never have to set up their first project alone. That's the whole shift, and it's why activation is finally a problem you can solve well instead of paper over.

What this is: a working model for getting a stranger to their first real win, fast, with the agent carrying the load. What it isn't: a list of onboarding tricks, or a pitch for adding more tooltips. We're going to treat activation the way a game designer treats Level 1, because games are the one discipline that takes total strangers from confused to capable with no manual, and they never sit anyone down for a lecture first.

The goal, in one line: get a new user to a result they'd tell a coworker about, on day one, with the agent doing the heavy lifting, so they come back tomorrow.



The model everything runs on

The game: Level Design for software

ACTIVATION IS LEVEL 1 · THE AGENT IS THE COMPANION · THE LOOP IS THE ENGINE

Before the tactics, the model the whole stack runs on. If you only take one idea from this protocol, take this one.

Think of your product as a game, and your user as a player who starts as a complete stranger and, if things go right, becomes a power user who can't imagine working without it. Getting them there isn't a phase that ends after week one. It's a series of levels they grow through, and the good ones never stop teaching as they go. Not Angry Birds. Think a deep RPG, where each level hands the player a skill they'll need for the next one, and new abilities unlock as they master the old ones. Mastery is the destination, not the exit.

That game has three levels, and they're the three protocols in this stack:

- **Level 1, Activation.** The first real win. A stranger reaches a result that makes them think "okay, this works." This protocol.
- **Level 2, Retention.** The win becomes a habit. They come back, get better, and the product opens up as they do.
- **Level 3, Expansion.** They hit the edges of what they've got, reach for more power, and start bringing other people in.

What moves a player from one moment to the next is a loop: a **cue** that pulls them in, an **action** they take, and a **reward** that pays off the action and sets up the next cue. Cue, action, reward, next cue. That loop is the engine of the whole game. It's how you get someone through activation, and it's the exact same mechanic that later carries them through retention and expansion. The first win you'll design in this protocol is just the first turn of that loop.

And in agent-led software, the game finally has a proper companion. In an old product the user had to read the manual and find their own way. With an agent in the product, you have a character who can greet the player, hand them the right first quest, do the hard parts for them, and point at the next move the moment they're ready. The agent is to your onboarding what a great game's design is to a player: it sequences the challenge so the user grows without ever feeling taught. Like product-led growth made the product the go-to-market, agent-led onboarding makes the agent the activation layer.

Everything that follows is how to design Level 1 so the agent can run it.

The whole protocol, one page

Read the rest for the why, or jump to the Activation Protocol on page 12.

The game

- Activation is Level 1: a stranger reaching their first real win, the same day they sign up.
- The agent runs the level. The user should never set up their first project alone.
- Onboarding isn't a separate flow bolted on top. It runs through the product's AI, in the same voice the user will keep using forever.

The win

- A win is the smallest result worth their time, not a finished setup. "Completed setup" is work; "found where you're losing sales" is a win.
- If you can't say it as a sentence, it isn't a win yet. It's a feature.
- Test it four ways: does it connect to their job, can they reach it solo in one sitting, would they tell a coworker, would it pull them back tomorrow.

Proof beats value

- Not every product can deliver real value in session one. Snippets, integrations, and imports take time.
- When you can't show value yet, show proof the value is coming: sample data, a quick grader, a live preview built from their world. Proof buys the time real value needs.

The path

- Three steps or fewer from signup to the win. Everything else moves to after the win or gets cut.
- Let the agent do the hard work: it builds the first project, fills the blanks, makes the boring decisions. The product produces something before it asks the user to.

The proof moment

- Name the win out loud, in their language, the moment it lands. Silent success reads as failure.
- Then hand them the next cue. The win answers one question and raises a better one; the agent points at it. That's the loop starting.



Why users leave before they start

THE COSTUME ISN'T THE PRODUCT

Most activation problems get misdiagnosed, and the wrong diagnosis is always the same: the product is too confusing, so we need more guidance on top of it. More tooltips, a longer tour, a setup checklist with a progress bar. So the team ships a layer of overlay guidance and waits for the activation curve to move. It doesn't, or barely.

The reason it doesn't is that overlay tools optimize for the wrong job. They let a growth team ship something that looks like onboarding without touching the thing that actually decides activation: the product itself. The first-run experience, the empty state, the default settings, what the product does before the user does anything. A tooltip pointing at a blank dashboard doesn't make the blank dashboard less blank. Most onboarding tools dress up the activation problem. They don't fix it.

The blank screen is the real culprit, and AI made a fancier version of it. The blank chat box is the new bad onboarding. A product that opens with "What can I do for you today?" and three generic example prompts you can't relate to hasn't removed onboarding, it's removed the floor. I've felt this myself testing new AI products: it looks great, I try two things, I never figure out what it's actually good at for me, and I drift away. Not from friction. From ambiguity. Nobody rage-quits these products. They just quietly stop opening them.

Agent-led onboarding is the fix, and it isn't "add a chatbot." A chatbot dropped on top of an unchanged product is plugin theater with a new face. The fix is that the agent is built into the product and runs the getting-started experience as its first real conversation with the user. It greets them, asks what they came for, and starts building. The user shouldn't be able to feel a seam between "the onboarding" and "the product." There's no separate onboarding mode to graduate out of; there's the agent, and it's the same companion on day one and in month six.

Stop asking how to walk the user through a confusing first screen, and start asking how the agent gets them to a win on that screen. The rest of this protocol answers the second question.



Define the same-day win

IF YOU CAN'T SAY IT AS A SENTENCE, IT'S NOT A WIN YET

You can't design Level 1 until you've decided what beating it means. So before anything else, define the win: the smallest meaningful result a new user can reach in their first session that makes them think the time was worth it. One result. Small on purpose. Not "see everything the product can do." One concrete moment: see where visitors drop off on a page, send one campaign, get one real answer about their business.

The most useful thing you can do here is force the win into a single sentence:

In their first session, a new **[user type]** will **[reach this win]** so they **[do their actual job]**.

Filled in: "In their first session, a new ecommerce founder will see which page is losing them the most sales, so they can fix it and make more money." That names who, the result, and the job. The sentence matters because it exposes the fakes. "Complete setup so they..." so they what? Setup isn't a win. It's the tax you charge before the win, and the agent's job is to get that tax as close to zero as possible.

Two traps the sentence fights off. The **feature trap**: "they'll create their first dashboard" is a feature; the win is what the dashboard tells them. Keep asking "so that?" until you hit something they wanted before they met your product. And the **everything trap**: making the first win big to make it impressive. Big wins need setup, setup needs time, and time is what a first-session user won't give you. A small real win beats a big theoretical one.

Then pressure-test the candidate against four questions. A real win passes all four, and the one it fails tells you what to fix.

TEST	THE QUESTION
Job	Connects to the job they're hiring the product for, or just a feature you like?
Session	One person, one sitting, no waiting on a teammate, import, or traffic? Most wins fail here.
Colleague	Would they describe it to a coworker, unprompted? Too small to mention is too small to matter.
Return	If they got only this, would they come back tomorrow?

The win is the target the whole level points at, and the agent can't drive toward a target you haven't named. Name it as a sentence, test it four ways, then build everything else backward from it.



Proof beats value

WHEN YOU CAN'T SHOW THE VALUE YET, SHOW PROOF IT'S COMING

Here's a tenet most activation advice skips, and it's where a lot of genuinely good products lose people on day one. Not every product can deliver real value in the first session. Some need data, time, or setup before any value can exist. Analytics needs traffic. A heatmap needs you to install a snippet and then wait for visitors. A CRM needs your deals in it. Pretending you have a same-day value win when the value is honestly days away is how you design a beautiful flow to a destination the user can't reach yet.

So split it. If the user can reach the win in session one with no waiting, that's a **VALUE WIN**, and chapters 4 through 6 are about clearing the path to it. If the value genuinely needs setup or time first, you're on a **PROOF PATH**, and your first-session job changes: you can't show them their result yet, so you show them believable proof the result is coming. Proof doesn't replace value. It buys you the time and effort real value needs, by earning enough belief that the user does the setup instead of leaving.

I live this at Heatmap, where I'm head of product. Heatmap can't show you your own data on day one, because your snippet isn't live and no traffic has come through yet. So we don't fake a value win. We do two things instead. We run a grader that scans the page you give us right now and hands back real suggestions based on what we actually see, so you get something useful before you've installed anything. And we let you load a sample store, so you can see exactly what your own heatmaps and funnels will look like, and learn the platform, before your data shows up. Neither one is your value. Both are proof, and proof is enough to get someone to install the snippet and come back.

This is the "show the solved cube" move from game design. Every Rubik's Cube ships solved; you scramble it yourself, which is why nobody believes it's impossible. Onboarding can do the same: show the user a finished version of where they're headed before they've earned it. A populated demo built from their industry. Another customer's end state. A preview of their own dashboard. It doesn't spoil the journey, it makes the journey feel survivable.

And this is where an agent earns its keep, because it can build the solved cube live and personal. "Here's a mock version of the dashboard you could have by Friday, built from a store like yours. Want me to start setting up the real one?" A generated, disposable, made-for-you preview is something a static product could never afford to ship. The agent can.

Figure out which path you're on before you design anything else, because it changes the whole first session. Most teams who say they "have no same-day win" actually have a proof problem they never named.



04

Clear the path, let the agent walk it

THE USER SHOULD NEVER SET UP THEIR FIRST PROJECT ALONE

You know the win. Now the route to it. Walk it yourself, from the signup button to the moment the win lands, and count the distinct steps as they actually exist today, not as you wish they were. The number is almost always too high. Sign up, verify, fill in a profile, pick a plan, take a tour, connect a source, configure a setting, and somewhere far down there, the win. Every step is a small cliff a share of people fall off. Stack enough and the win becomes theoretical: it's in there, but almost nobody reaches it.


Aim for three steps or fewer from signup to the first win. You won't always hit it, but the target forces the real question: which of these steps truly has to happen before the win, and which are just there because someone added them and nobody removed them. Three moves get you there.

- **Cut or defer everything that isn't load-bearing.** Most "setup" can wait until after the win, once the user has a reason to care. Win first, housekeeping second.
- **Do the hard work for them.** Sample data so the screen isn't empty. Templates so they're editing, not creating from nothing. Smart defaults so the boring decisions are already made. The product produces something before it asks the user to produce anything. A blank screen that says "get started" is the product handing its job to the user.
- **Make the one action obvious.** One clear next action per screen, obvious in two seconds without reading. First-time users scan and click; they don't read.

And in an agent-led product, the agent does that work in the open. It doesn't point at the empty project and explain how to fill it. It builds the first project with them and narrates as it goes: "I connected your store, pulled in your top five pages, and set up the first funnel. Take a look." The user learns the product by watching the agent operate it on their own stuff, which is worth more than any tour. This is the move great games make with their characters: the companion does it for you on quest one, with you on quest two, and watches you do it on quest three. Quest one is right now.

One more thing the agent should hold: two goals at once. Serve the request in front of it, and quietly steer the user one level up. Answer the question, then point at the island they can't see yet. "Done. While that import finishes, want me to set up your first alert?" The user comes away feeling like they figured the product out themselves, which is exactly the feeling you want.

A short path with the agent carrying the load beats a long path with a helpful tooltip on every step. Cut the path to the bone, then let the agent walk what's left.



05

The scripted first quest

A WIN ON SAMPLE DATA IS A DEMO. A WIN ON THEIR DATA IS A STORY.

This is where activation borrows most directly from how games teach, because the moves translate almost one to one, and the agent is what makes them possible.

- **Run it on their stuff.** Tetris grips you with no story because everything in it is happening to you, not to a character. A first win on a sample dataset is something the user watched happen to someone else. A first win on their actual account, even a tiny one, is something that happened to them, and they remember it. Get to their data fast. Superhuman's onboarding ran on your real inbox from the first minute, and that's most of why people who used it still talk about it. An agent can earn this in minutes: connect one source, produce one real thing from it, narrate what just happened.
- **Plant a question, don't push a tour.** A tour answers questions nobody asked. The stronger move is to make the user ask one of their own, about their own business: what would this say about my funnel? An agent does it in one line: "I noticed your signups spike on Tuesdays. Want to know why?" Curiosity about their own numbers is the best onboarding engine there is, and it costs the agent one sentence.
- **Be the candy clerk, not the blank box.** Faced with sixty flavors, people freeze; faced with a blank "what do you want to do," they freeze harder. The fix isn't fewer options, it's a short, opinionated menu that still feels open: "There's a lot you can do here, but for someone in your role, most people start with these three." You give them the feeling of freedom and a real place to start.

The thread through all three: the agent does the steering, so the user gets to feel like they figured it out themselves. That feeling is the win behind the win.

Script the first quest so the agent runs it on the user's real world, sparks their curiosity, and hands them a confident place to start. That's how a stranger gets their first result without reading a thing.



Land the win, start the loop

SILENT SUCCESS READS AS FAILURE. SAY THE WIN OUT LOUD.

The user reaches the win. Don't let the moment pass quietly. A result the product doesn't acknowledge is a result the user might not even notice, and a win they didn't notice didn't happen, as far as activation is concerned. Two parts to landing it.

Name the win, in their language. Not "dashboard created." Not "task complete." Say what it means to them: "Your first report's ready, and it looks like 23% of your budget is going to the channel that converts the worst." The generic confirmation is the product talking about itself; the specific one is the product talking about their business, and only one of those gets remembered. Think of a basketball net. It adds nothing to the rules of the game; it exists so everyone can see, instantly, that the ball went in. Your success moment is the net. In an agent-led product the agent is perfectly placed to call it, because it just did the work and can report the result like a colleague: "Done. Connected your store, and the one thing that jumps out is your checkout page loses 40% of people."

Then hand them the next cue. Right after the win, point at one clear next thing, not a menu. The win answered a question and raised a better one; the next move is the agent handing them the thread to pull. "You found the leak on your pricing page. Want to see the three pages losing you the most overall?" That single prompt is the loop turning over: the reward from the first win becomes the cue for the next action.

And that's the real reason the win matters. The first win isn't the finish line of activation, it's the first turn of the loop that runs the entire game. Get a stranger to one real result, name it, and hand them the next cue, and you haven't just activated them. You've started the engine that carries them into Level 2, retention, where the habit forms, and eventually Level 3, expansion. Because the onboarding lives in the agent, there's no graduation and no handoff. The same companion that got them their first win is the one that levels them up for as long as they use the product.

Activation's job is done the moment the user has a real win and a clear reason to come back tomorrow. Land that, and you've stopped the biggest leak most products have.



The worksheet

The Activation Protocol

DO IT YOURSELF, OR HAND IT TO YOUR FAVORITE LLM

Work through it for your product, or hand it to your favorite LLM along with this protocol and your product details and have it draft each answer with you, then sharpen. Either way you finish with a defined, tested, reachable same-day win and a first session the agent can run.

STEP 1: WRITE YOUR WIN SENTENCE

In their first session, a new _____ (user type)
will _____ (the smallest real win)
so they _____ (their actual job).

STEP 2: RUN THE FOUR TESTS

A fail means fix the win before you build anything else.

TEST	THE QUESTION	PASS / FAIL	
Job	Connects to their job, not just a feature?	☐	☐
Session	Reachable solo, one sitting, no waiting?	☐	☐
Colleague	Would they describe it to a coworker?	☐	☐
Return	Would it pull them back tomorrow?	☐	☐

STEP 3: VALUE OR PROOF

- ☐ **Value win: reachable in session one, no waiting. Go to Step 4.**
- ☐ **Proof path: value needs data, time, or setup first. Design the proof:**

Before their own data is ready, they'll see [the believable proof: sample data, a grader, a live preview] so they trust the value is coming and do the setup.

STEP 4: MAP AND CUT THE PATH

List every step from signup to the win as it works today, in order. Mark each: **K** keep · **D** defer · **C** cut.

Step count after cutting: (target: 3 or fewer)

STEP 5: HAND THE WORK TO THE AGENT

What the agent builds or fills in, so the user doesn't start from blank:

The first quest the agent runs on the user's real data:

The question the agent plants to spark curiosity:

STEP 6: WRITE THE WIN MOMENT

The success message, in their language (name the result, not the feature):

The next cue the agent offers right after the win:

All six filled in is your Level 1 design: a real win, proven reachable, on a path short enough to walk, run by the agent on the user's own world, ending in a moment they notice and a reason to come back. That's activation. Retention is the next level.



The dual use

Run this with your favorite LLM

THE PROTOCOL IS THE EXPERT; YOUR AI ASSISTANT DOES THE LOCALIZING

You can read this protocol and do the work yourself. The faster move, and the one it's built for, is to hand the whole thing to your favorite LLM and design your activation with you. Attach this protocol, give it a few paragraphs about your product (what it is, who it's for, what happens after signup today), then work through these asks in order.

ASK YOUR LLM	WHAT YOU'RE GETTING	CH.
"Based on my product, draft three candidate same-day wins as win sentences."	A shortlist of real first-session wins, phrased the right way.	02
"Run all three through the four win tests and tell me which survives."	The strongest win, and why the others fail.	02
"Can my product deliver this value in session one, or am I on a proof path? If proof, design the proof moment."	A clear value-or-proof call, and a sample-data or preview plan if you need one.	03
"Map the current path from signup to this win, then tell me what to cut or defer to get under three steps."	A shorter, sharper path with the dead weight flagged.	04
"Write the script for how my product's agent runs the first quest on the user's real data, including the curiosity question it plants."	The agent's first-session behavior, written out.	05
"Write the success message in the user's language, and the next cue the agent offers right after the win."	The win moment and the first turn of the loop.	06

Push back on your AI assistant the way you would on a teammate. If a win sentence reads like a feature, say so and make it try again. The model gives you the structure and the judgment; your taste decides what's actually right for your users. That last bit is the part that's yours.

**Get a stranger to one real win, name it,
and hand them the next move.
Everything else is in service of that.**

The Onboarding Loop Protocols

BAL SIEBER · ACTIVATION, LEVEL 01